

Seeds: Random effect logistic regression

Ported to PyMC by Abraham Flaxman and Kyle Foreman, 2/15/2012 from <http://www.openbugs.info/Examples/Seeds.html>

This example is taken from Table 3 of Crowder (1978), and concerns the proportion of seeds that germinated on each of 21 plates arranged according to a 2 by 2 factorial layout by seed and type of root extract. The data are shown below, where r_i and n_i are the number of germinated and the total number of seeds on the i -th plate, $i = 1, \dots, N$. These data are also analysed by, for example, Breslow and Clayton (1993).

The model is essentially a random effects logistic, allowing for over-dispersion. If p_i is the probability of germination on the i -th plate, we assume

$$r_i \sim \text{Binomial}(p_i, n_i)$$

$$\text{logit}(p_i) = \alpha_0 + \alpha_1 x_{1,i} + \alpha_2 x_{2,i} + \alpha_{1,2} x_{1,i} x_{2,i} + b_i$$

$$b_i \sim \text{Normal}(0, \tau)$$

where $x_{1,i}, x_{2,i}$ are the seed type and root extract of the i -th plate, and an interaction term $\alpha_{1,2} x_{1,i} x_{2,i}$ is included.

$\alpha_0, \alpha_1, \alpha_2, \alpha_{1,2}, \tau$ are given independent weakly informative priors.

```
In [1]: import pylab as pl
import pymc as mc
```

```
In [2]: ### data
# germinated seeds
r = pl.array([10, 23, 23, 26, 17, 5, 53, 55, 32, 46, 10, 8, 10, 8, 23, 0, 3, 22, 15,
# total seeds
n = pl.array([39, 62, 81, 51, 39, 6, 74, 72, 51, 79, 13, 16, 30, 28, 45, 4, 12, 41, 3
# seed type
x1 = pl.array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1])
# root type
x2 = pl.array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1])
# number of plates
N = x1.shape[0]
```

Here is the model in BUGS:

```
model
{
  for( i in 1 : N ) {
    r[i] ~ dbin(p[i],n[i])
    b[i] ~ dnorm(0.0,tau)
    logit(p[i]) <- alpha0 + alpha1 * x1[i] + alpha2 * x2[i] +
      alpha12 * x1[i] * x2[i] + b[i]
  }
  alpha0 ~ dnorm(0.0,1.0E-6)
  alpha1 ~ dnorm(0.0,1.0E-6)
  alpha2 ~ dnorm(0.0,1.0E-6)
  alpha12 ~ dnorm(0.0,1.0E-6)
  tau ~ dgamma(0.001,0.001)
  sigma <- 1 / sqrt(tau)
}
```

```

In [3]: ### hyperpriors
tau = mc.Gamma('tau', 1.e-3, 1.e-3, value=10.)
sigma = mc.Lambda('sigma', lambda tau=tau: tau**-.5)

### parameters
# fixed effects
alpha_0 = mc.Normal('alpha_0', 0., 1e-6, value=0.)
alpha_1 = mc.Normal('alpha_1', 0., 1e-6, value=0.)
alpha_2 = mc.Normal('alpha_2', 0., 1e-6, value=0.)
alpha_12 = mc.Normal('alpha_12', 0., 1e-6, value=0.)

# random effect
b = mc.Normal('b', 0., tau, value=np.zeros(N))

# expected parameter
logit_p = (alpha_0 + alpha_1*x1 + alpha_2*x2 + alpha_12*x1*x2 + b)

### likelihood
@mc.observed
def obs(value=r, n=n, logit_p=logit_p):
    return mc.binomial_like(r, n, mc.invlogit(logit_p))

```

BUGS uses Gibbs steps automatically, so it only takes 10000 steps of MCMC after a 1000 step burn in for this model in their example.

PyMC only uses Gibbs steps if you set them up yourself, and it uses Metropolis steps by default. So 10000 steps go by quickly, but the chain has not converged to the stationary distribution in this time, so predictions based on the samples will be inaccurate.

```

In [4]: m = mc.MCMC([alpha_0, alpha_1, alpha_2, alpha_12, b, tau, sigma, logit_p, obs])

# not long enough to mix
%time m.sample(10000, 1000, progress_bar=False)

```

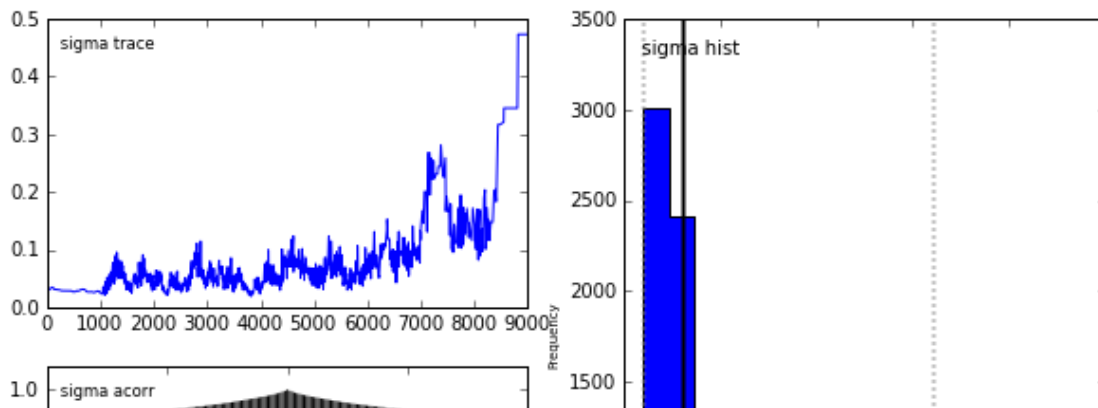
CPU times: user 6.75 s, sys: 0.11 s, total: 6.86 s
Wall time: 8.12 s

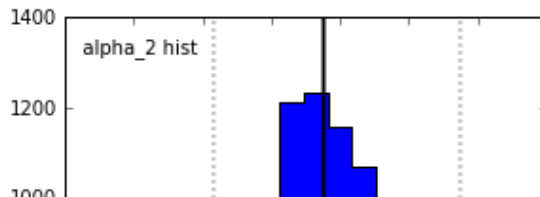
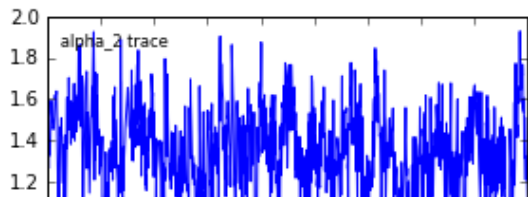
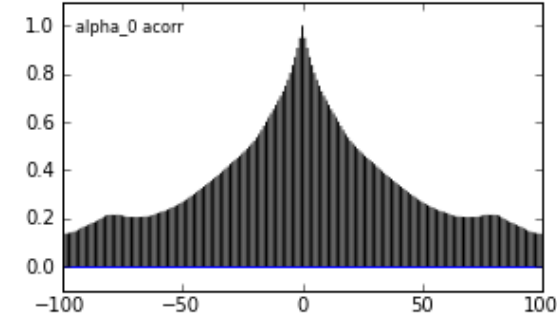
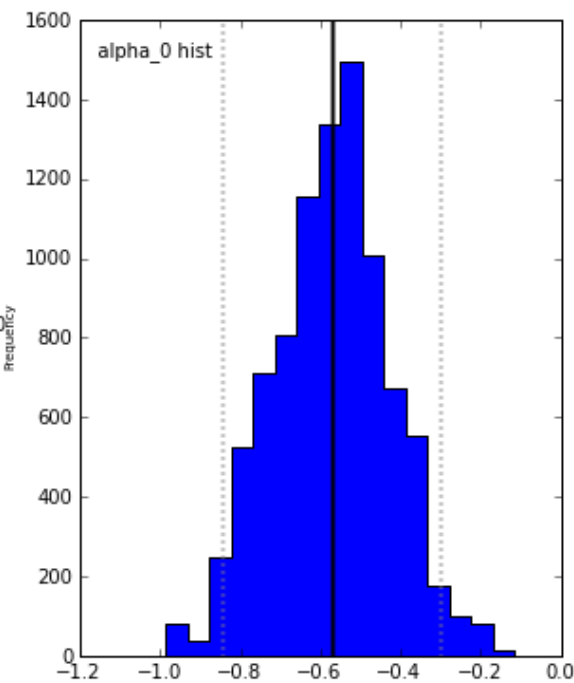
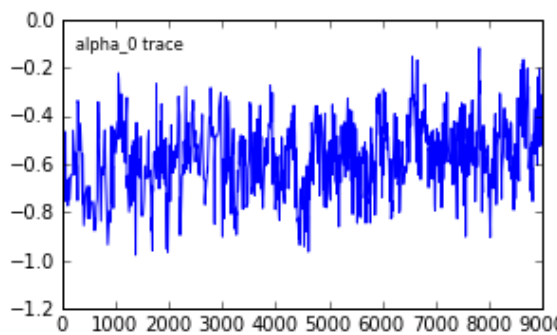
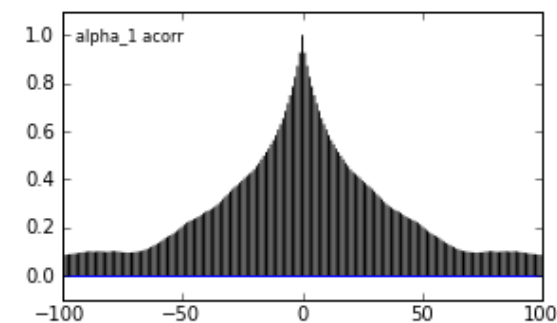
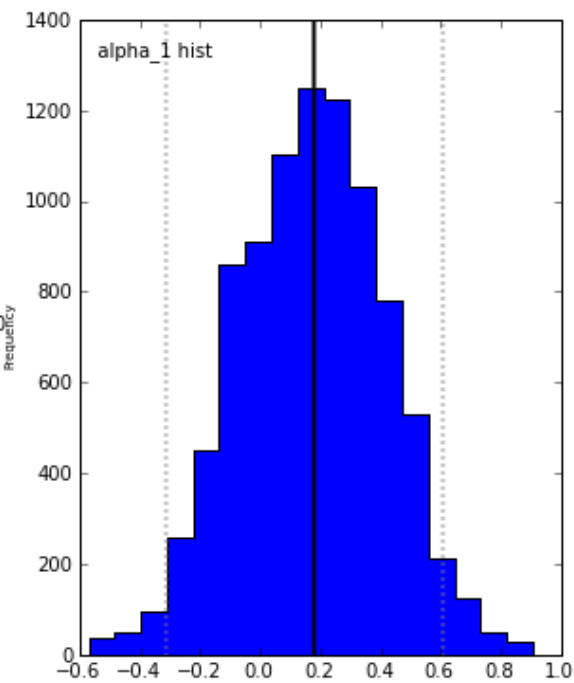
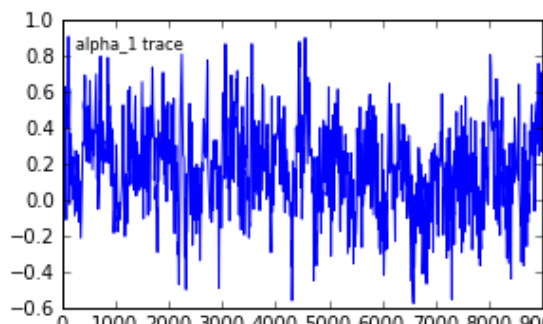
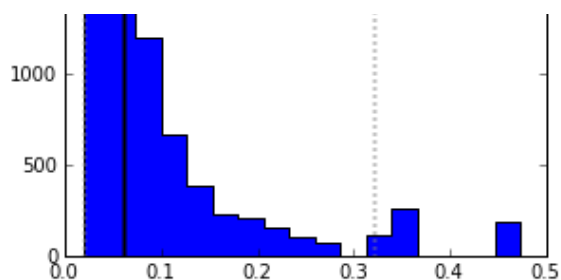
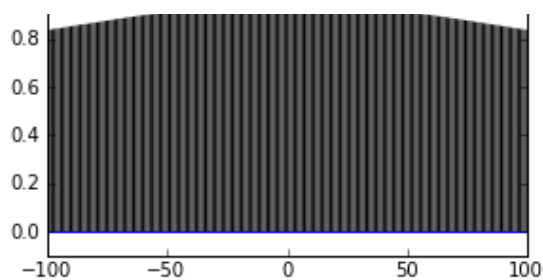
```

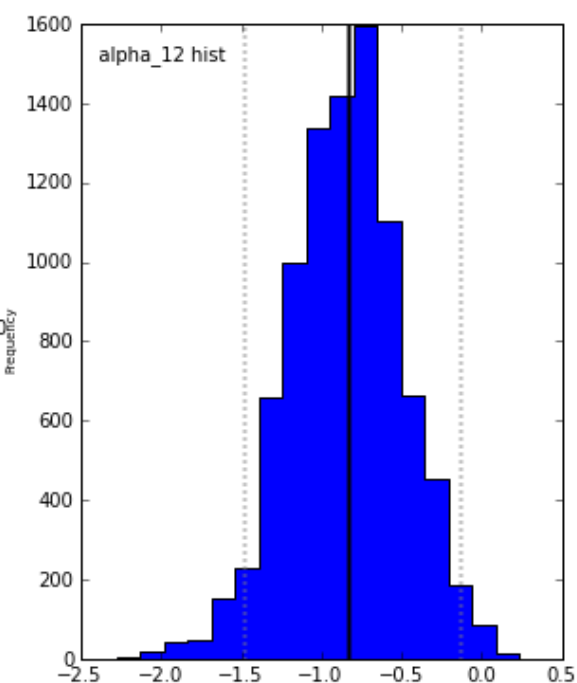
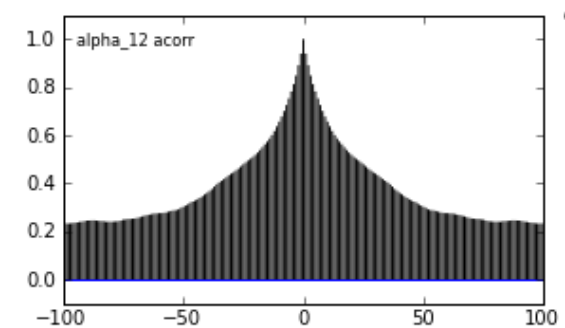
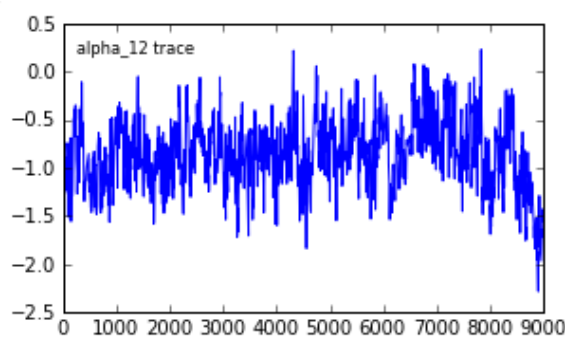
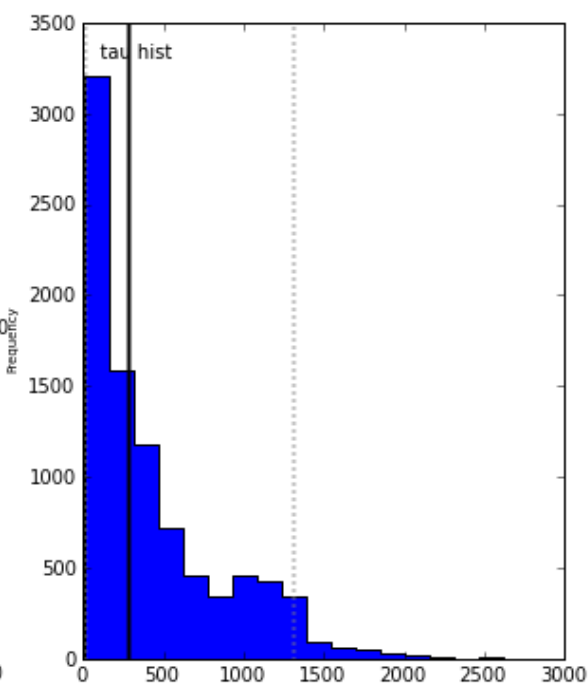
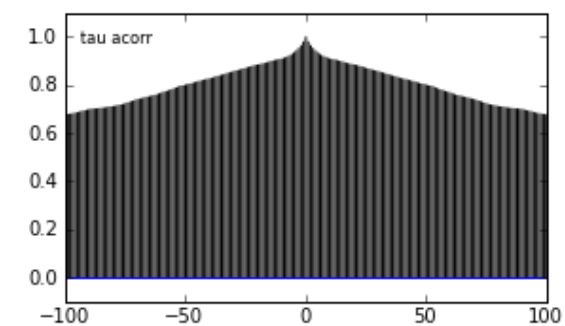
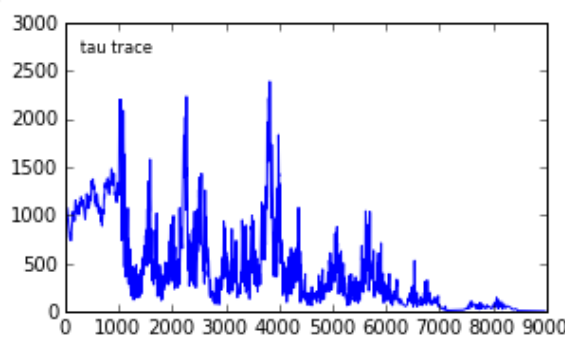
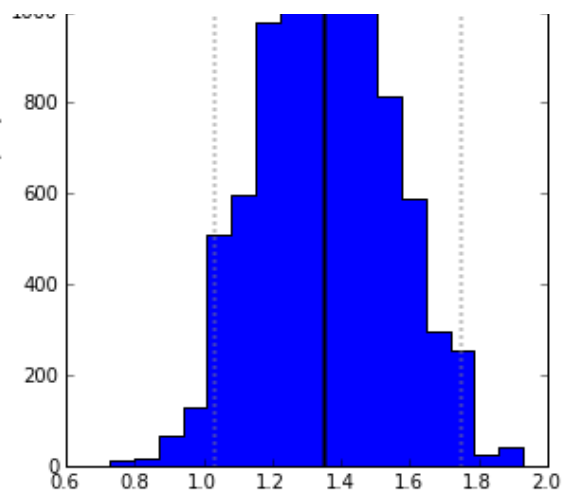
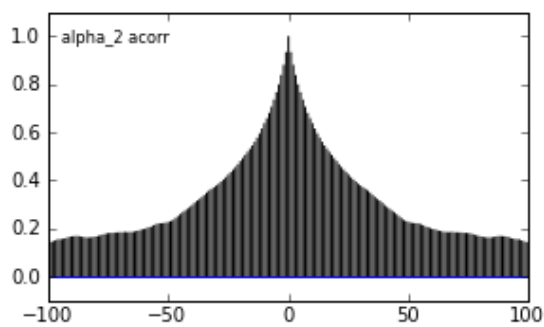
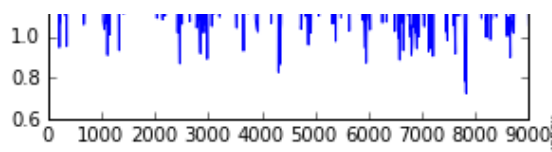
In [5]: mc.Matplot.plot(m)

```

Plotting sigma
Plotting alpha_1
Plotting alpha_0
Plotting alpha_2
Plotting tau
Plotting alpha_12







PyMC has a very sophisticated "Adaptive Metropolis" method, which adjusts the proposal distribution for the Metropolis step adaptively. Using this and running for longer provide more reliable results.

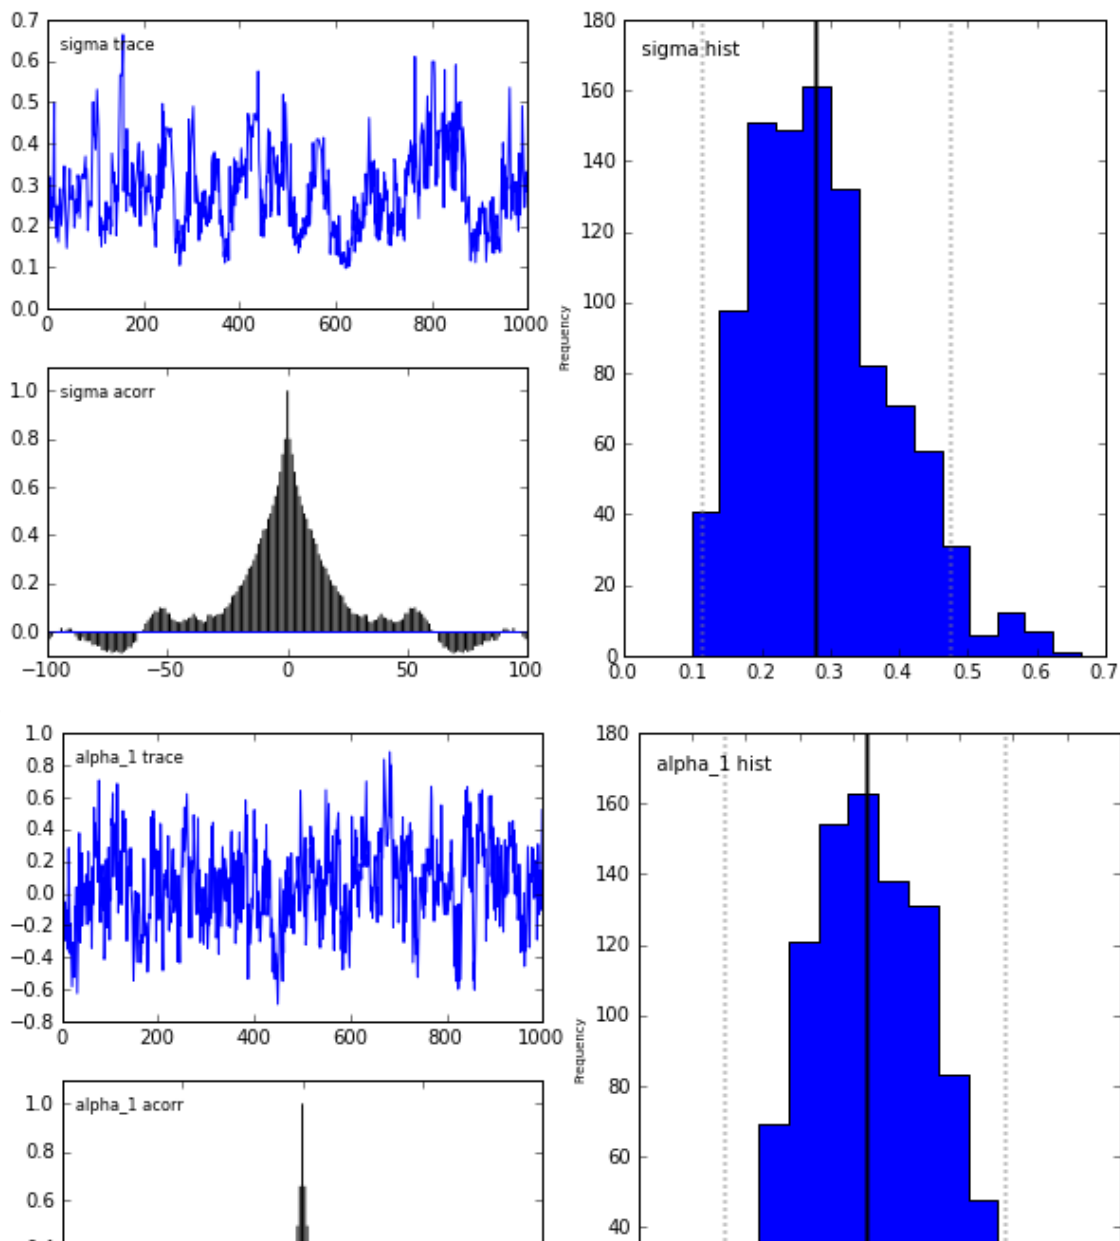
```
In [6]: m = mc.MCMC([alpha_0, alpha_1, alpha_2, alpha_12, b, tau, sigma, logit_p, obs])
m.use_step_method(mc.AdaptiveMetropolis, b)

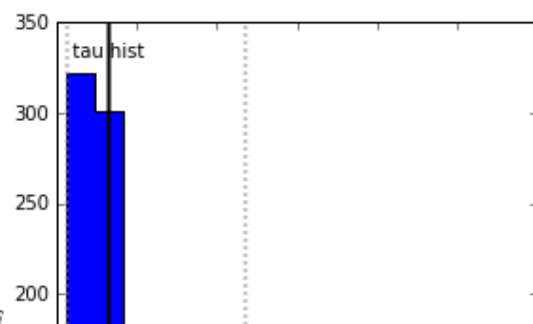
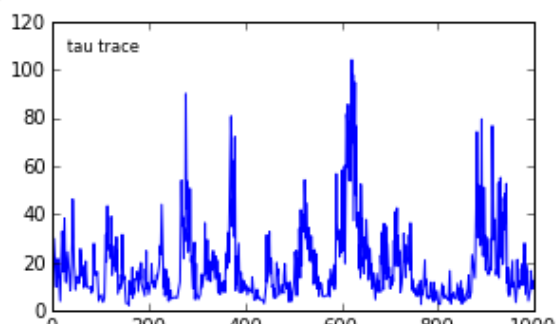
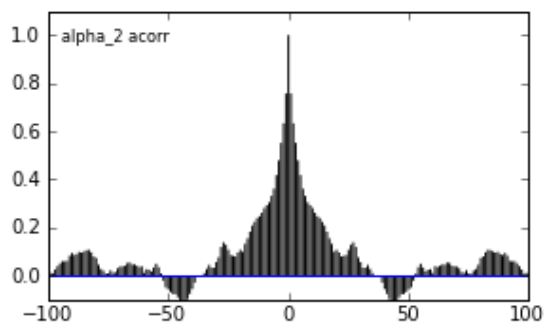
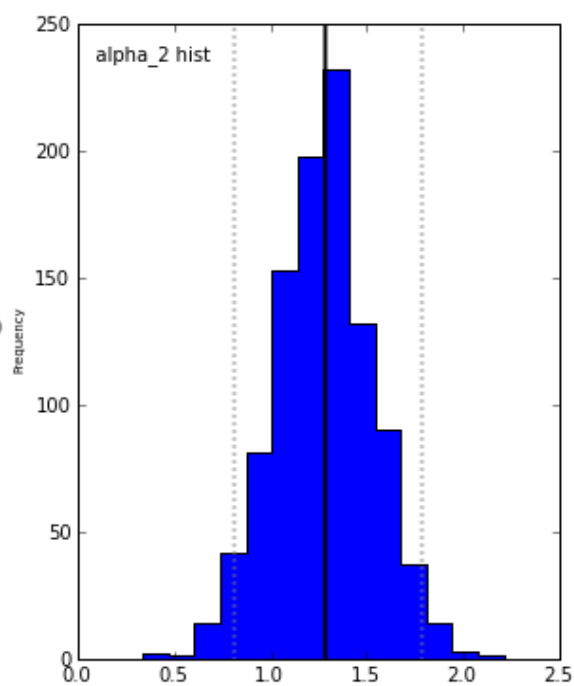
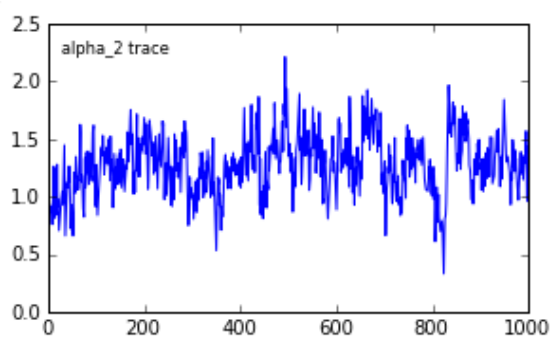
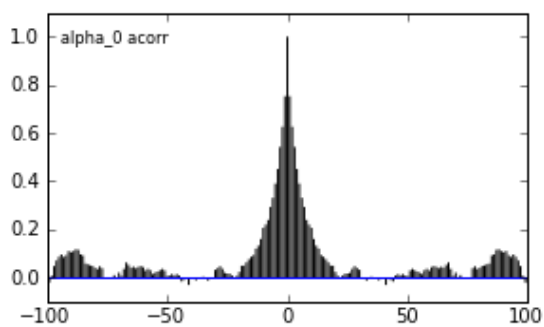
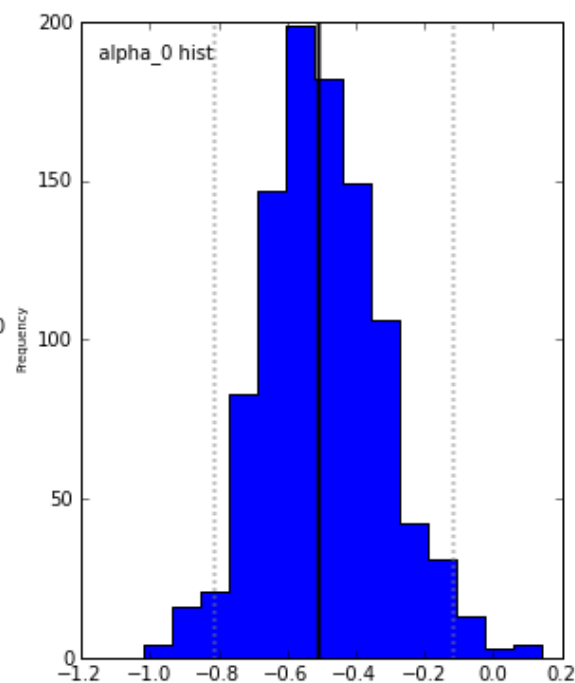
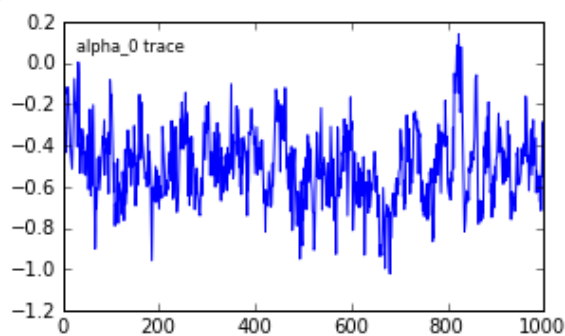
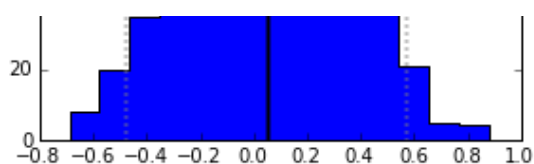
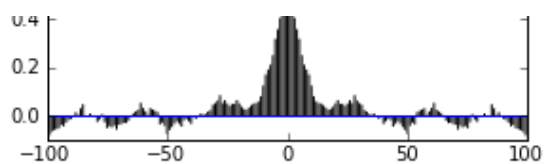
# about long enough, but maybe better initial values would help more
%time m.sample(20000, 10000, 10, progress_bar=False)
```

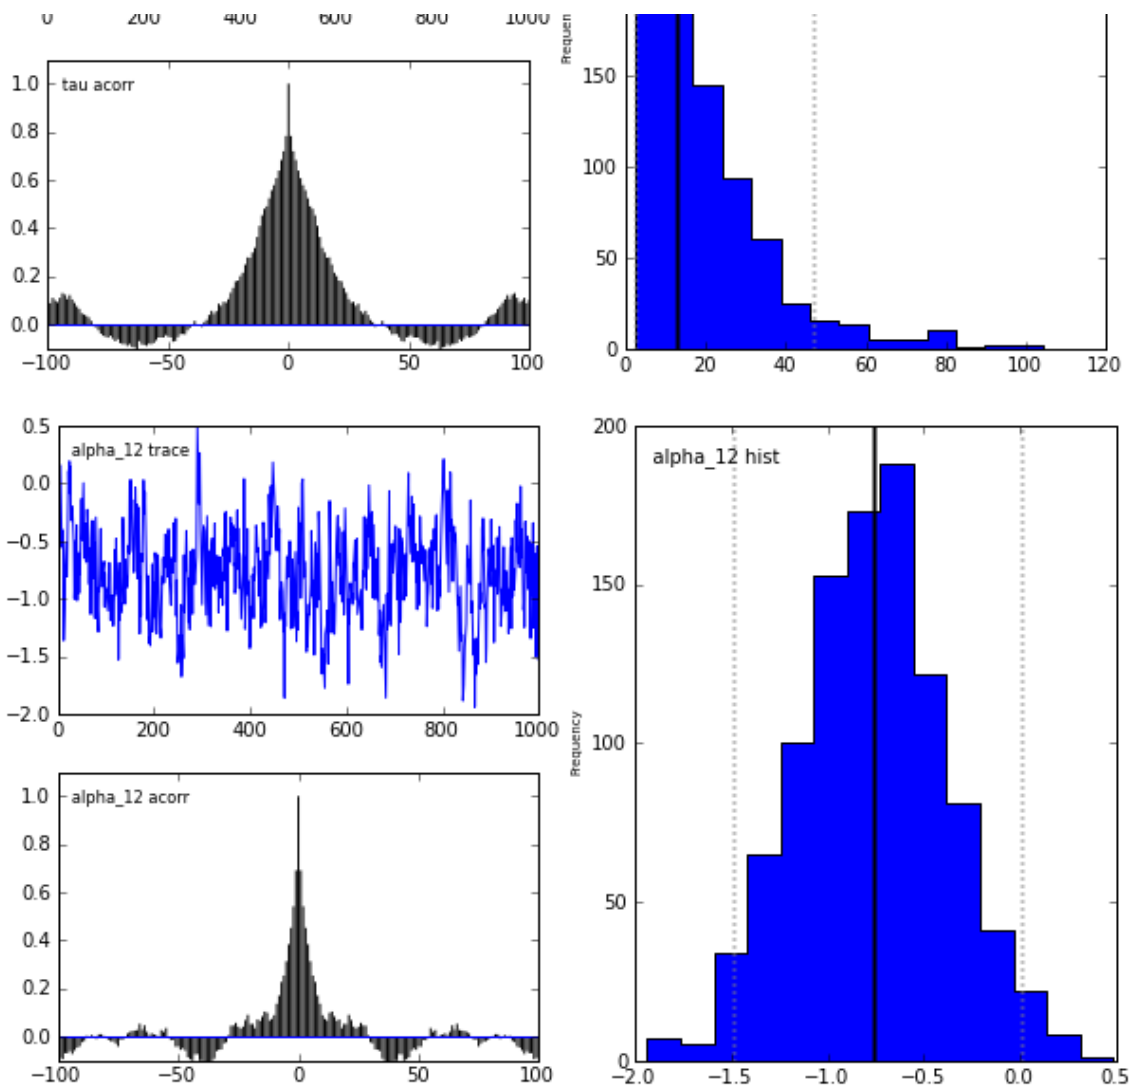
CPU times: user 13.35 s, sys: 0.33 s, total: 13.68 s
Wall time: 20.01 s

```
In [7]: mc.Matplot.plot(m)
```

Plotting sigma
Plotting alpha_1
Plotting alpha_0
Plotting alpha_2
Plotting tau
Plotting alpha_12







Starting from carefully chosen initial values and running the chain for even longer yields results that are completely mixed. This is going to give the same answer every time (approximately).

```
In [8]: tau.value=1.
map = mc.MAP([alpha_0, alpha_1, alpha_2, alpha_12, b, logit_p, obs])
map.fit(method='fmin_powell')

m = mc.MCMC([alpha_0, alpha_1, alpha_2, alpha_12, b, tau, sigma, logit_p, obs])
m.use_step_method(mc.AdaptiveMetropolis, b)

# a little longer, good initial values, but not Adaptive Metropolis. Does not converge
%time m.sample(200000, 100000, 100, progress_bar=False)
```

CPU times: user 125.07 s, sys: 2.74 s, total: 127.80 s

Wall time: 151.96 s

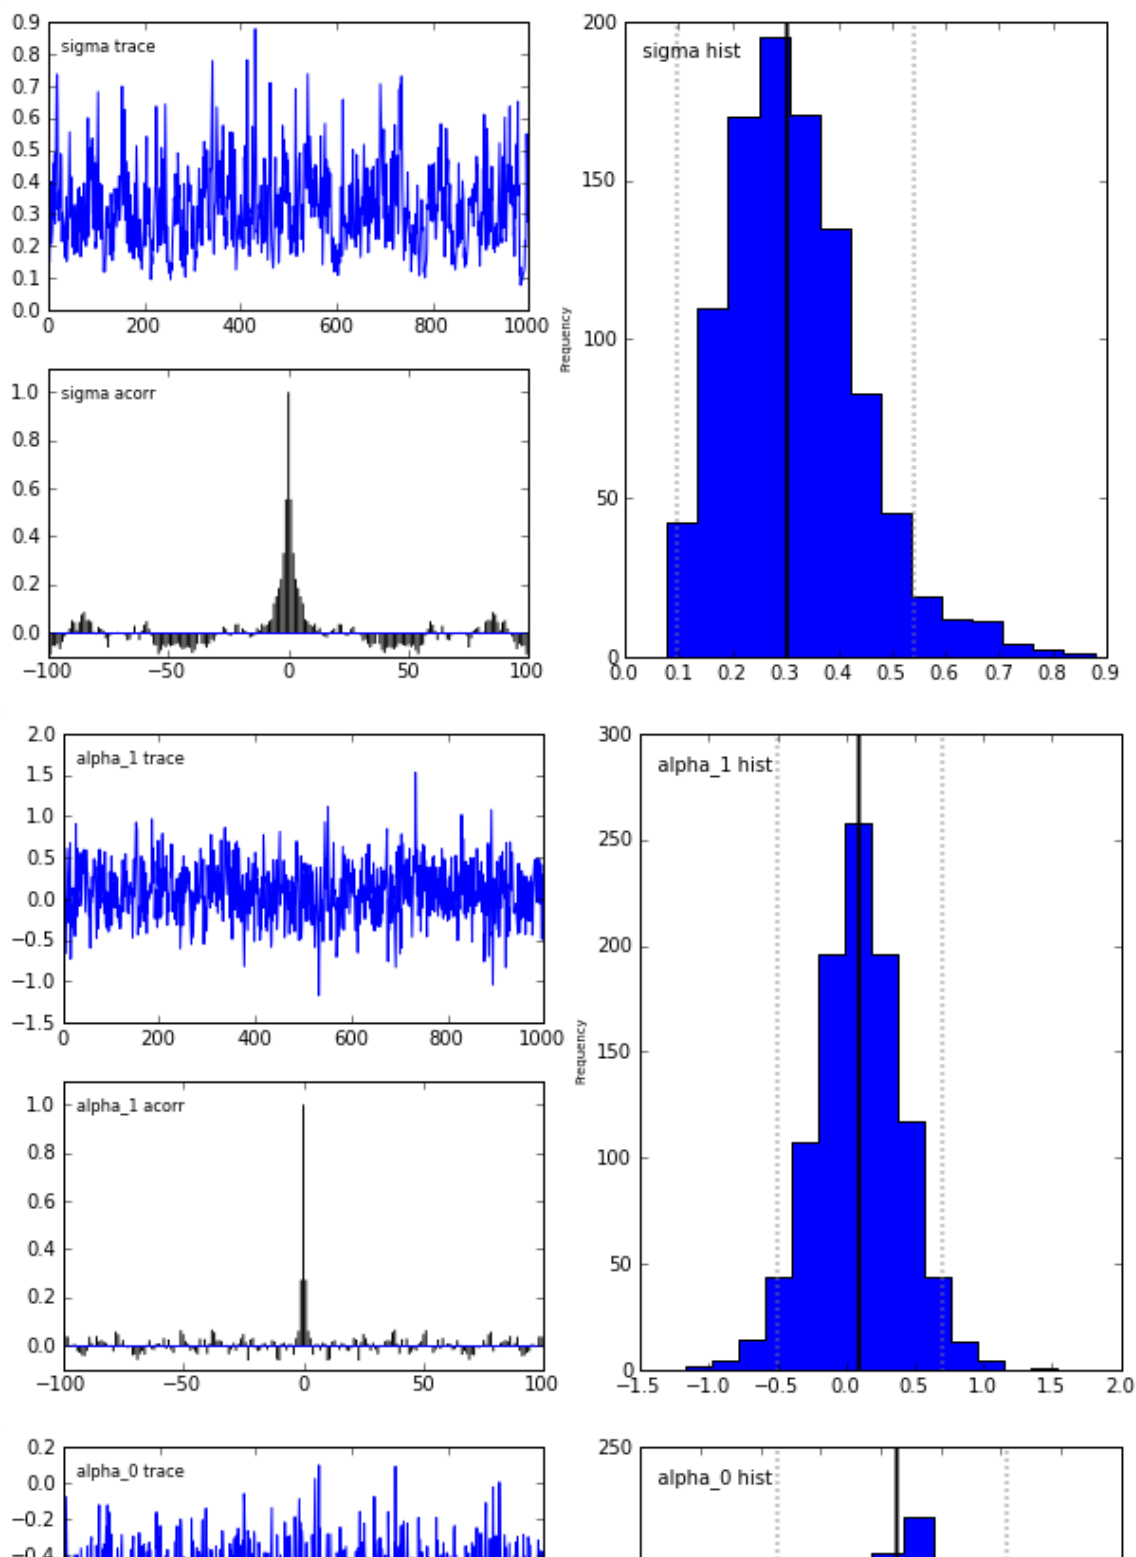
/home/abie/projects/env_pymc_dev/lib/python2.6/site-packages/pymc/StepMethods.py:1171: UserWarning:

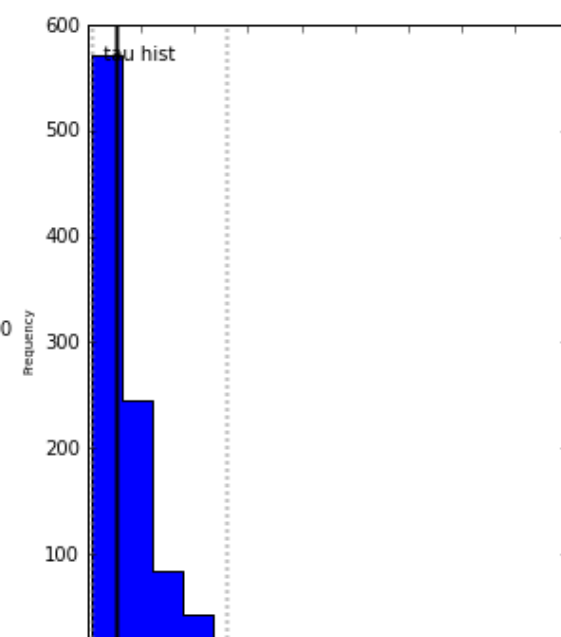
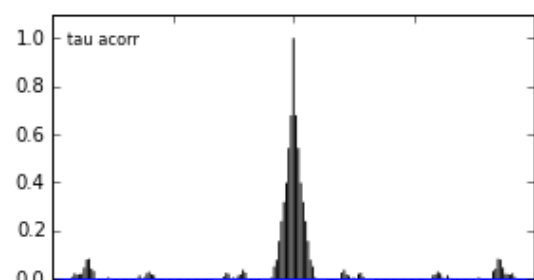
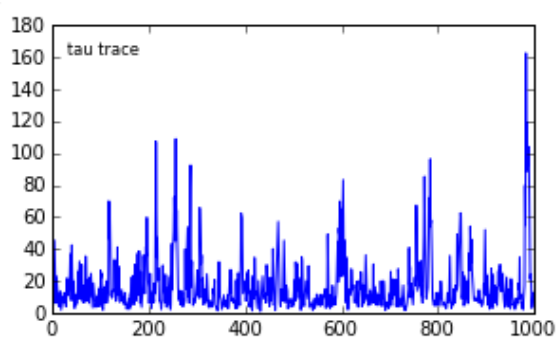
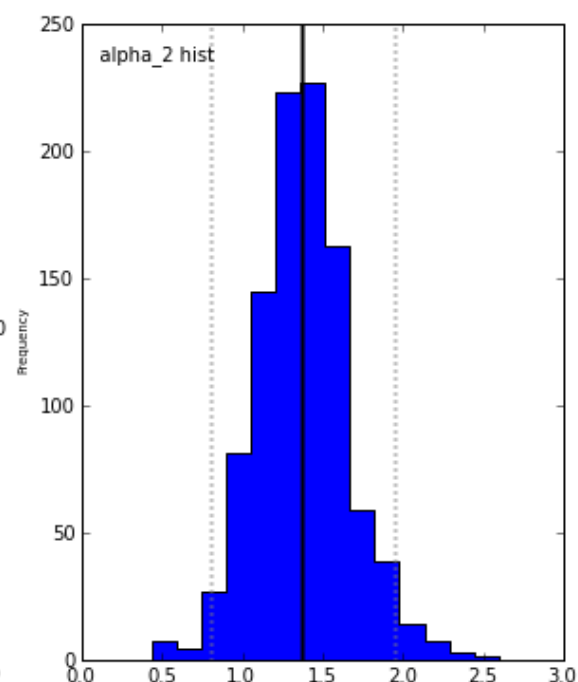
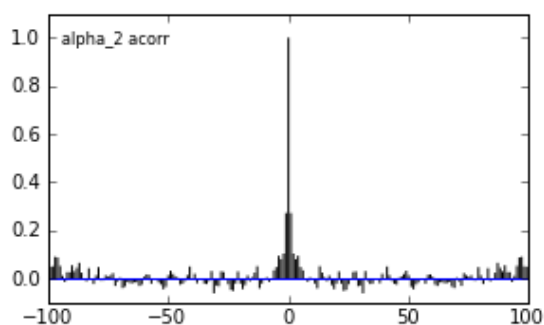
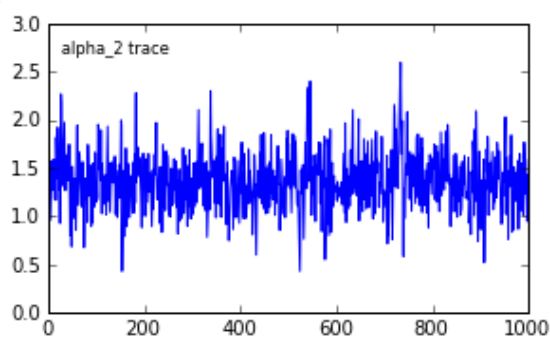
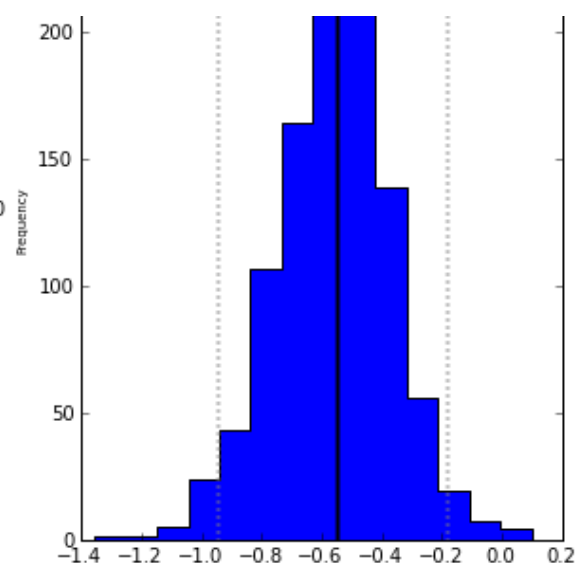
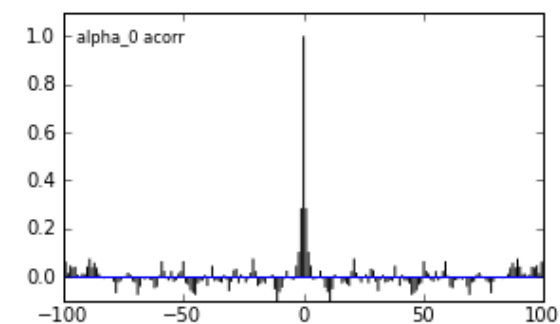
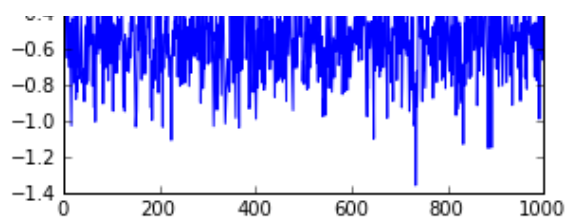
Covariance was not positive definite and proposal_sd cannot be computed by Cholesky decomposition. The next jumps will be based on the last valid covariance matrix. This situation may have arisen because no jumps were accepted during the last `interval`. One solution is to increase the interval, or specify an initial covariance matrix with a smaller variance. For this simulation, each time a similar error occurs, proposal_sd will be reduced by a factor .9 to reduce the

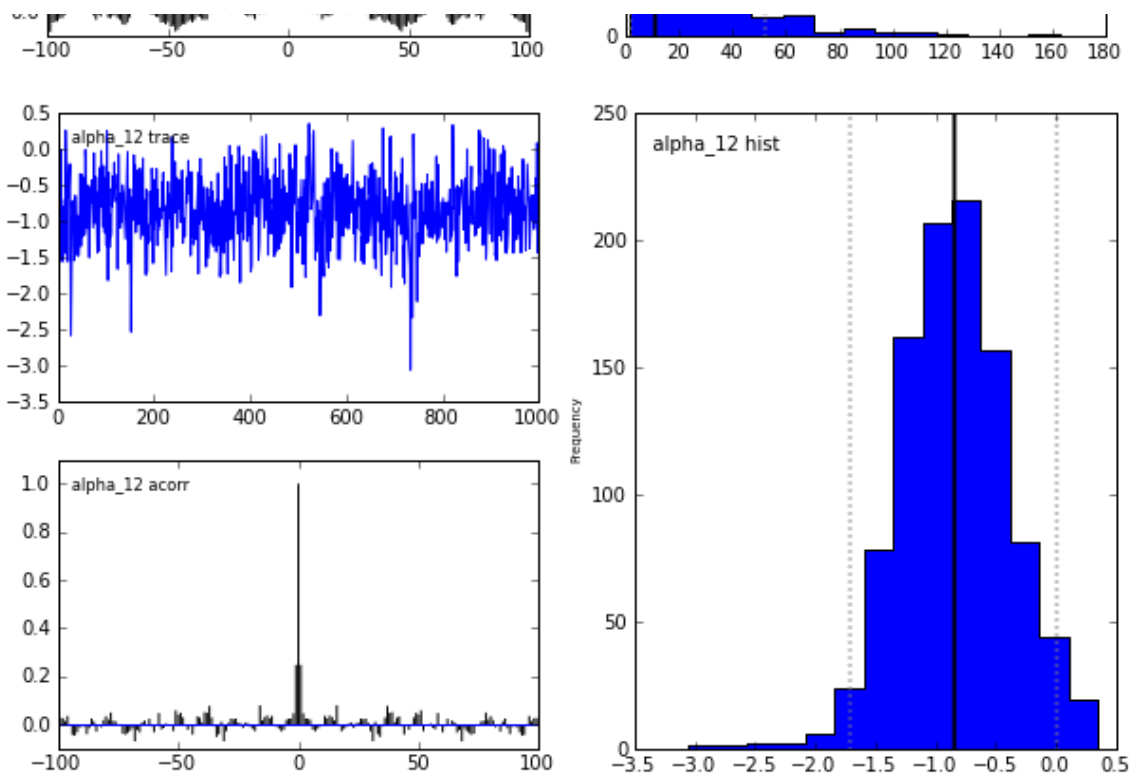
jumps and increase the likelihood of accepted jumps.
warnings.warn(adjustmentwarning)

```
In [9]: mc.Matplot.plot(m)
```

Plotting sigma
Plotting alpha_1
Plotting alpha_0
Plotting alpha_2
Plotting tau
Plotting alpha_12







Results

BUGS results:

A burn in of 1000 updates followed by a further 10000 updates gave the following parameter estimates:

	mean	sd
alpha_0	-0.55	0.19
alpha_1	0.08	0.30
alpha_12	-0.82	0.41
alpha_2	1.35	0.26
sigma	0.27	0.15

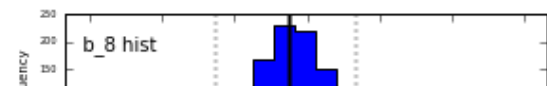
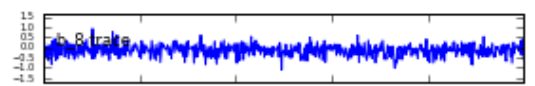
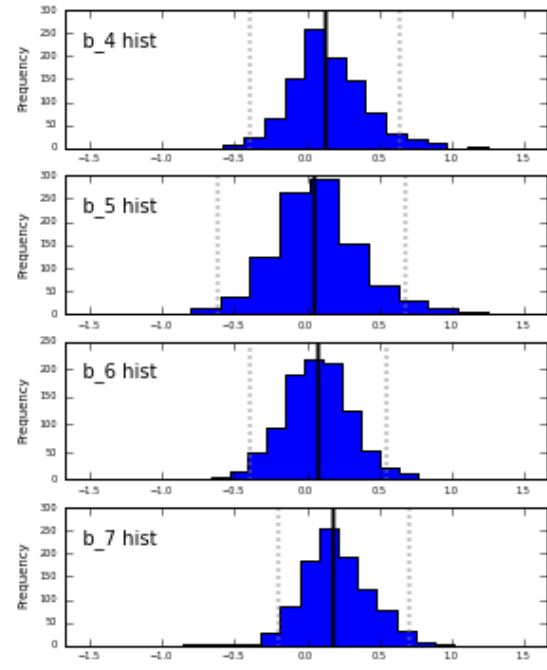
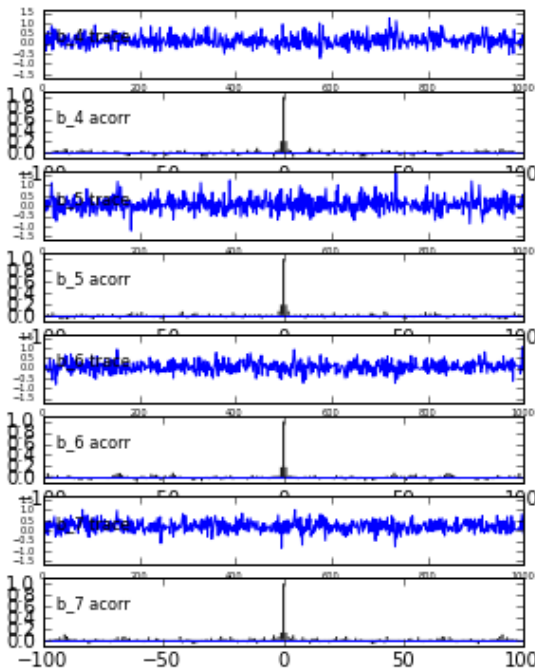
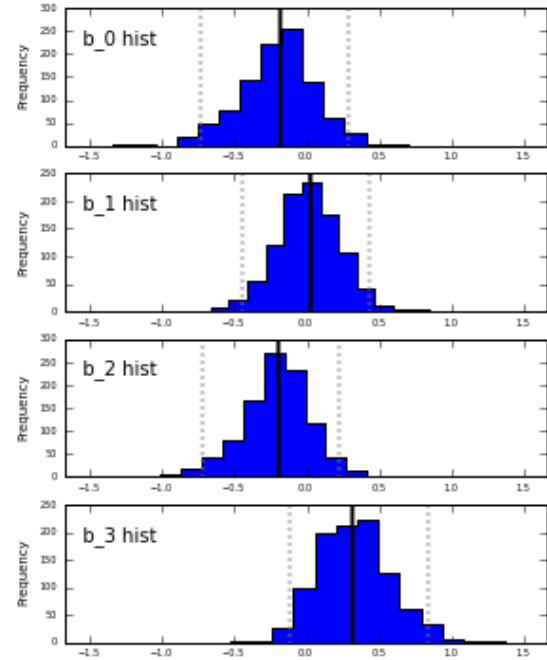
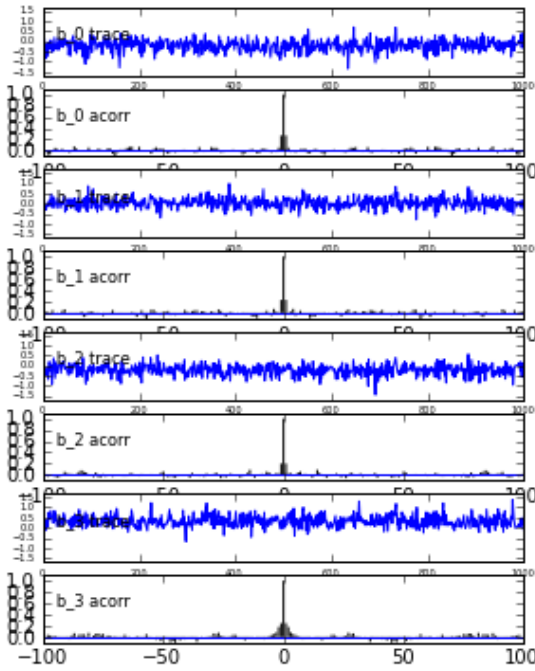
```
In [10]: for node in [alpha_0, alpha_1, alpha_12, alpha_2, sigma]:
          stats = node.stats()
          print '%10s\t%1.2f \t%.2f' % (node.__name__, stats['mean'], stats['standard deviation'])
```

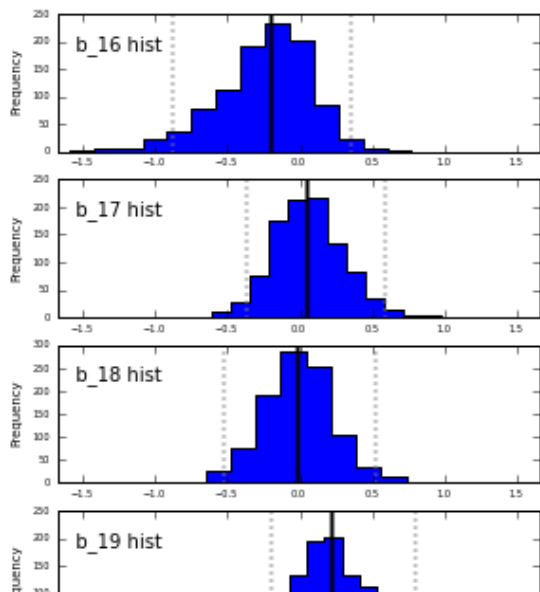
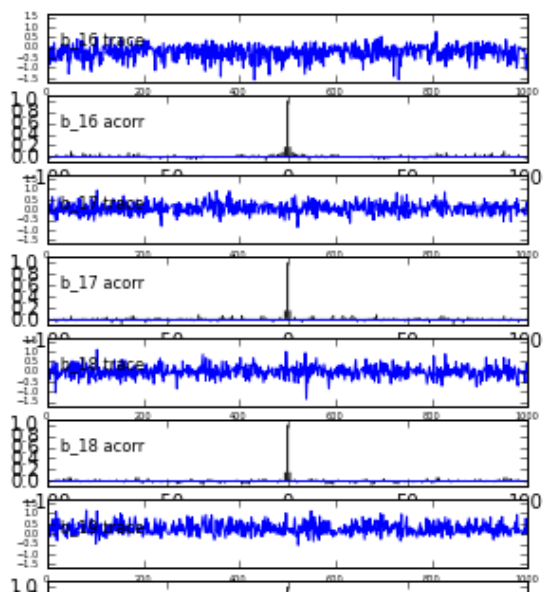
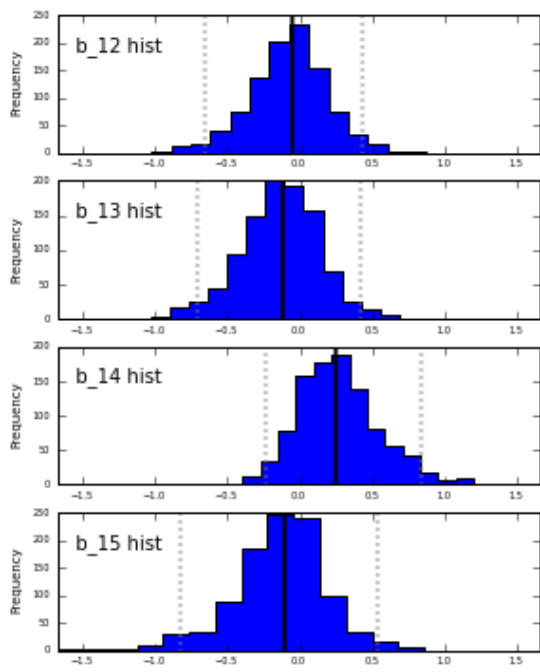
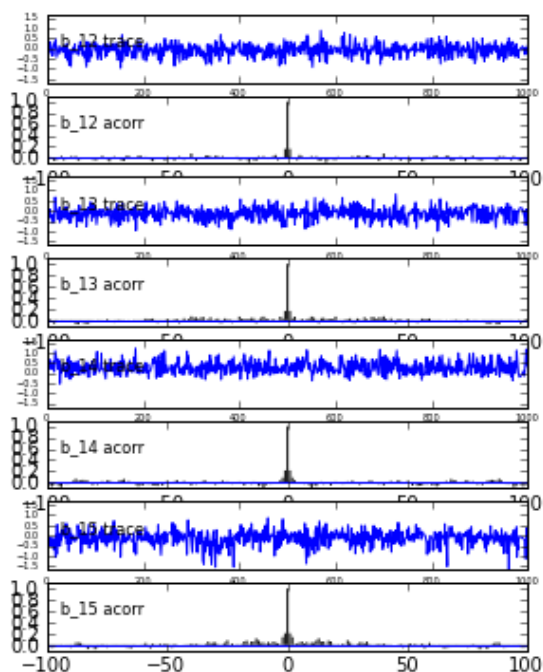
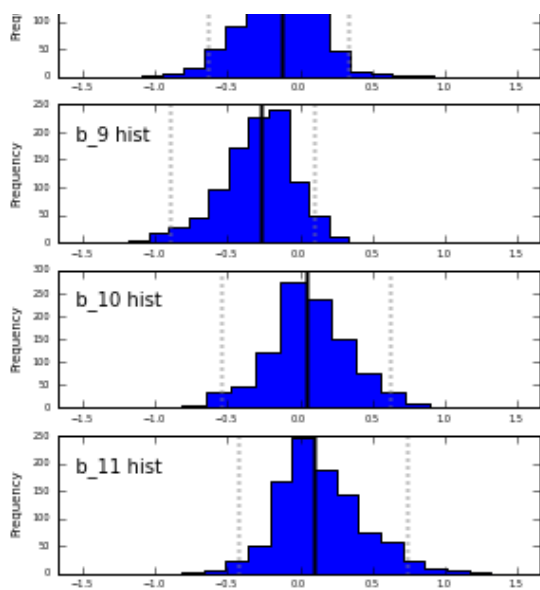
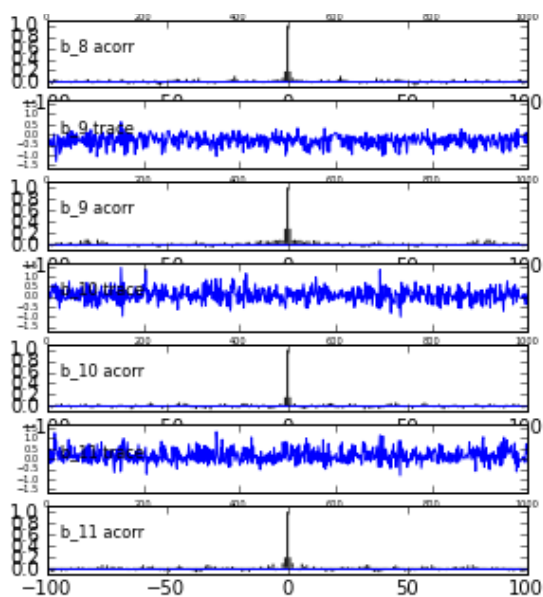
alpha_0	-0.56	0.19
alpha_1	0.08	0.32
alpha_12	-0.85	0.45
alpha_2	1.37	0.29
sigma	0.32	0.12

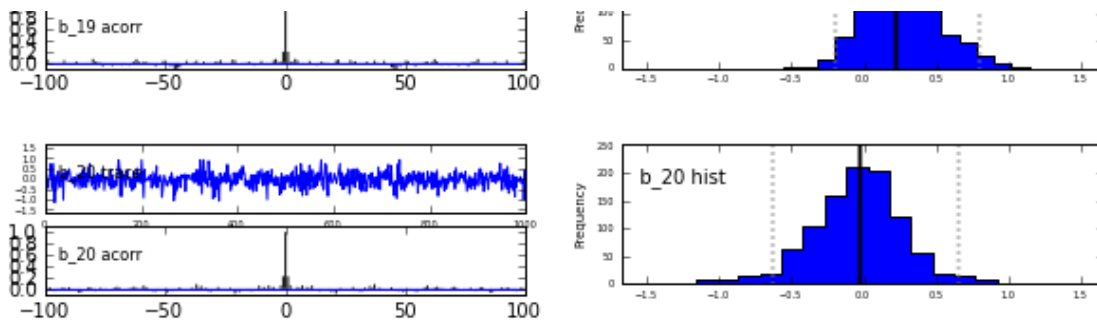
```
In [11]: mc.Matplot.plot(b)
```

```
Plotting b_0
Plotting b_1
Plotting b_2
Plotting b_3
Plotting b_4
Plotting b_5
Plotting b_6
Plotting b_7
```

Plotting b_8
 Plotting b_9
 Plotting b_10
 Plotting b_11
 Plotting b_12
 Plotting b_13
 Plotting b_14
 Plotting b_15
 Plotting b_16
 Plotting b_17
 Plotting b_18
 Plotting b_19
 Plotting b_20







Further exploration

Now that I've seen that it can work, how necessary is the Adaptive Metropolis? Necessary.

```
In [12]: tau.value=1.
map = mc.MAP([alpha_0, alpha_1, alpha_2, alpha_12, b, logit_p, obs])
map.fit(method='fmin_powell')

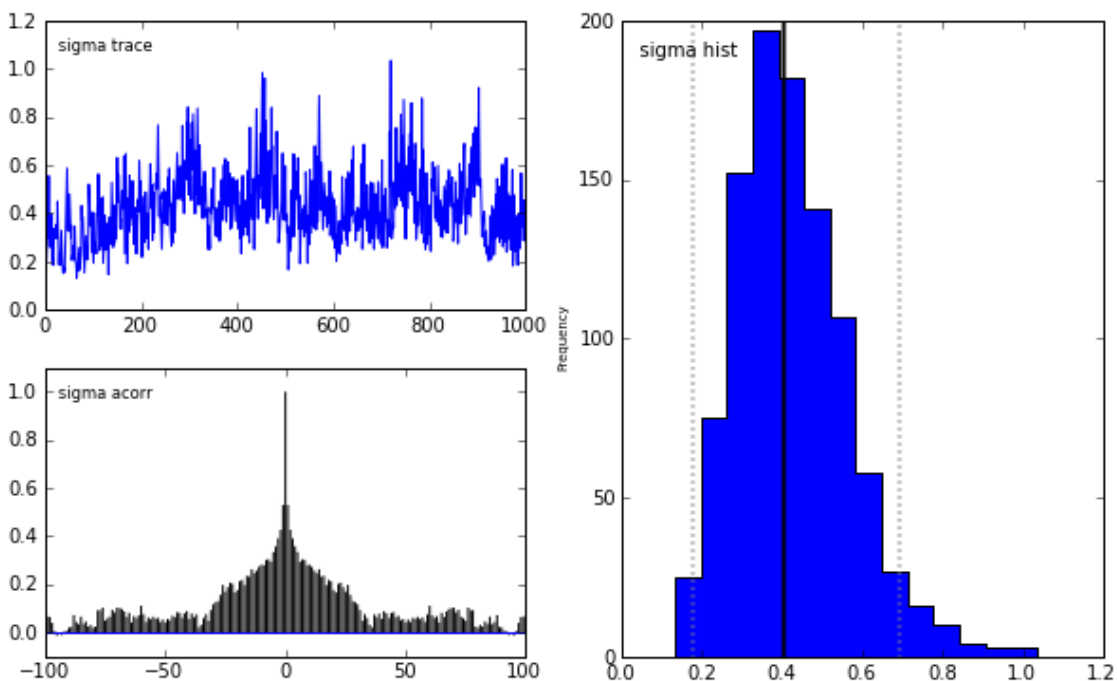
m = mc.MCMC([alpha_0, alpha_1, alpha_2, alpha_12, b, tau, sigma, logit_p, obs])

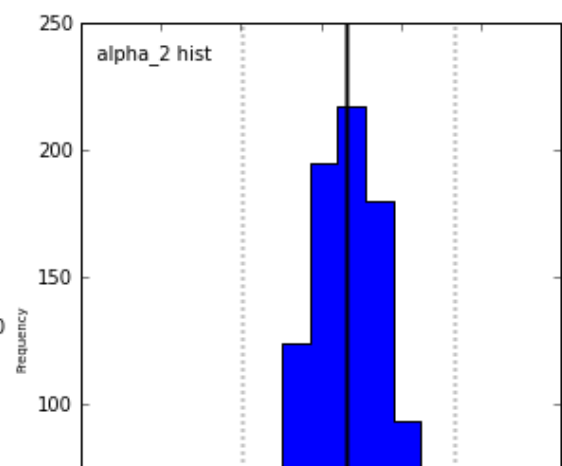
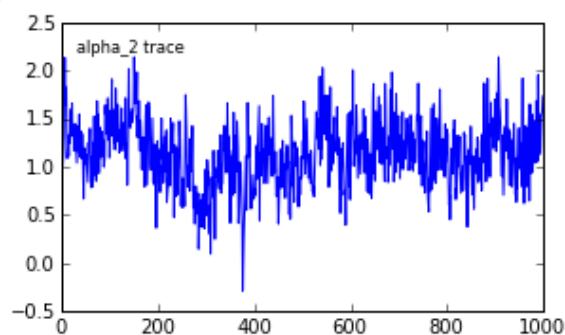
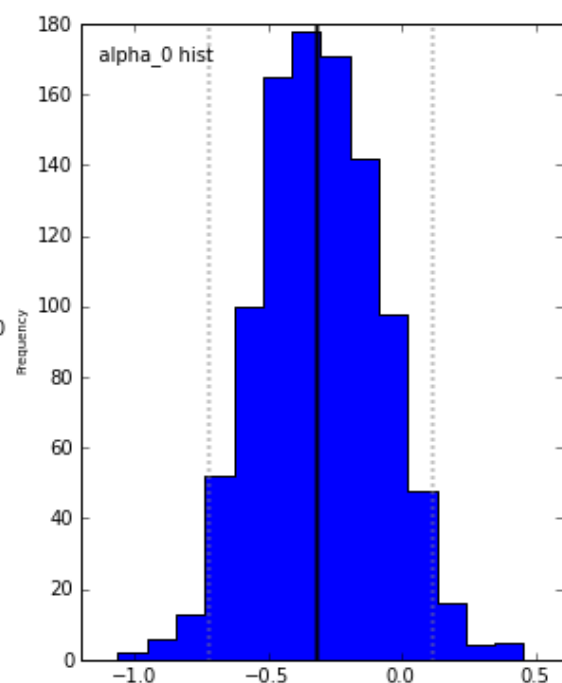
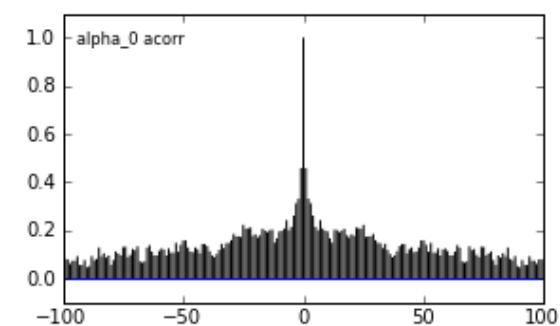
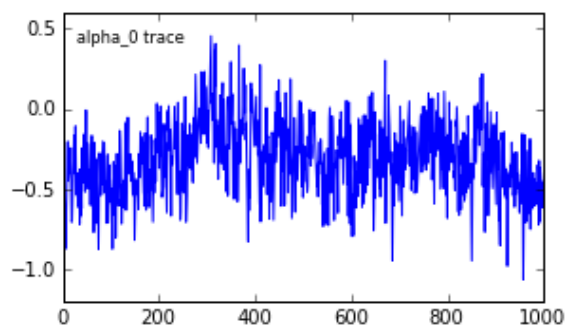
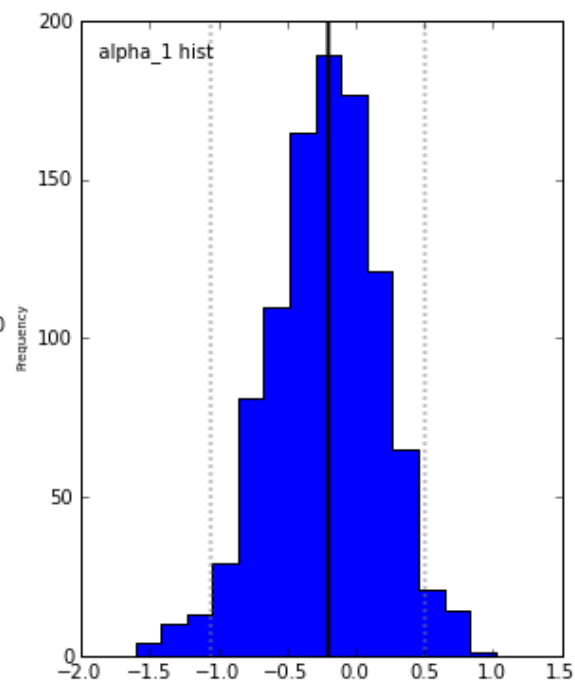
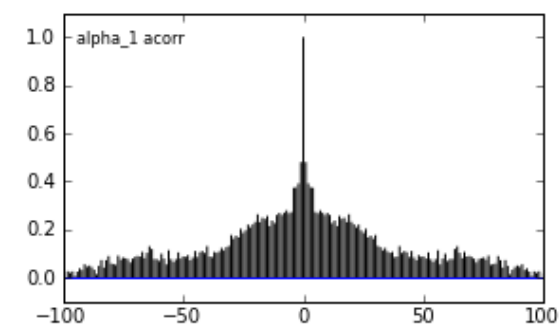
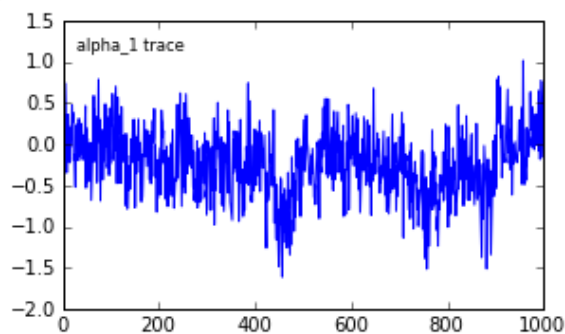
%time m.sample(200000, 100000, 100, progress_bar=False)
```

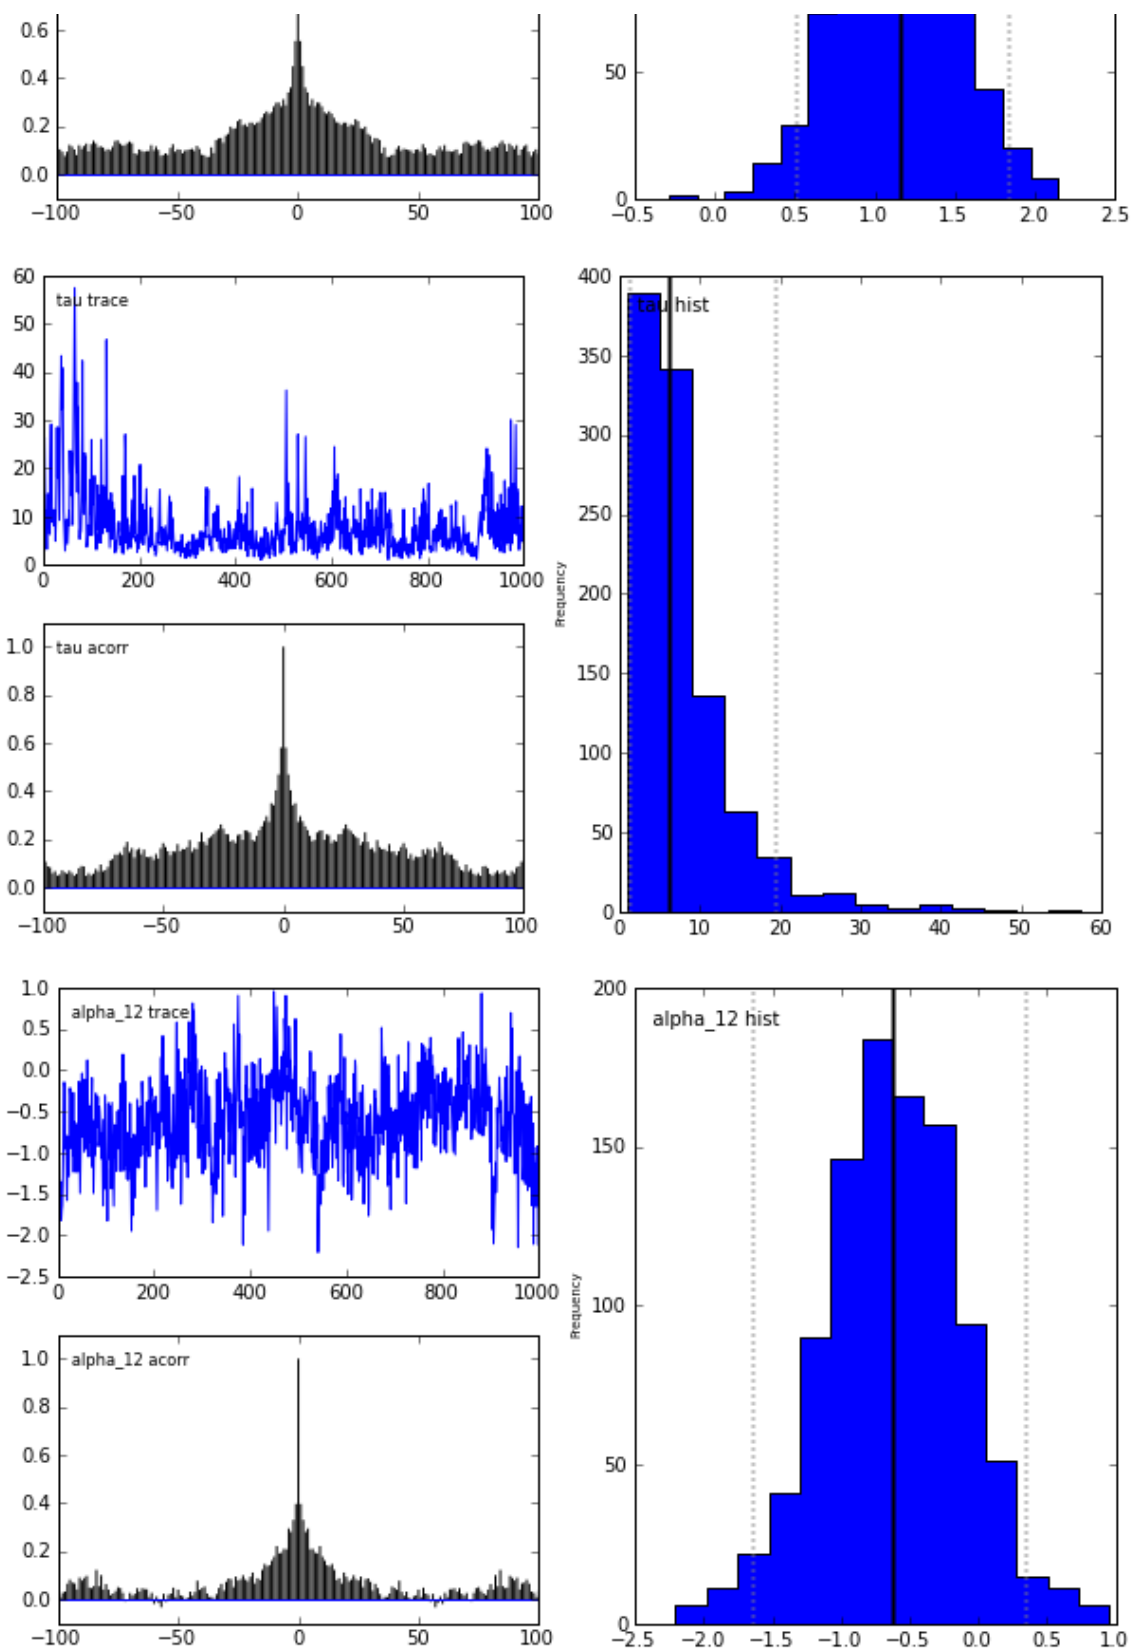
CPU times: user 122.45 s, sys: 0.70 s, total: 123.15 s
Wall time: 127.73 s

```
In [13]: mc.Matplot.plot(m)
```

Plotting sigma
Plotting alpha_1
Plotting alpha_0
Plotting alpha_2
Plotting tau
Plotting alpha_12







And how necessary are the carefully chosen initial values?

```
In [14]: ### hyperpriors
tau = mc.Gamma('tau', 1.e-3, 1.e-3, value=10.)
sigma = mc.Lambda('sigma', lambda tau=tau: tau**-.5)

### parameters
```

```

# fixed effects
alpha_0 = mc.Normal('alpha_0', 0., 1e-6, value=0.)
alpha_1 = mc.Normal('alpha_1', 0., 1e-6, value=0.)
alpha_2 = mc.Normal('alpha_2', 0., 1e-6, value=0.)
alpha_12 = mc.Normal('alpha_12', 0., 1e-6, value=0.)

# random effect
b = mc.Normal('b', 0., tau, value=np.zeros(N))

# expected parameter
logit_p = (alpha_0 + alpha_1*x1 + alpha_2*x2 + alpha_12*x1*x2 + b)

### likelihood
@mc.observed
def obs(value=r, n=n, logit_p=logit_p):
    return mc.binomial_like(r, n, mc.invlogit(logit_p))

m = mc.MCMC([alpha_0, alpha_1, alpha_2, alpha_12, b, tau, sigma, logit_p, obs])
m.use_step_method(mc.AdaptiveMetropolis, b)

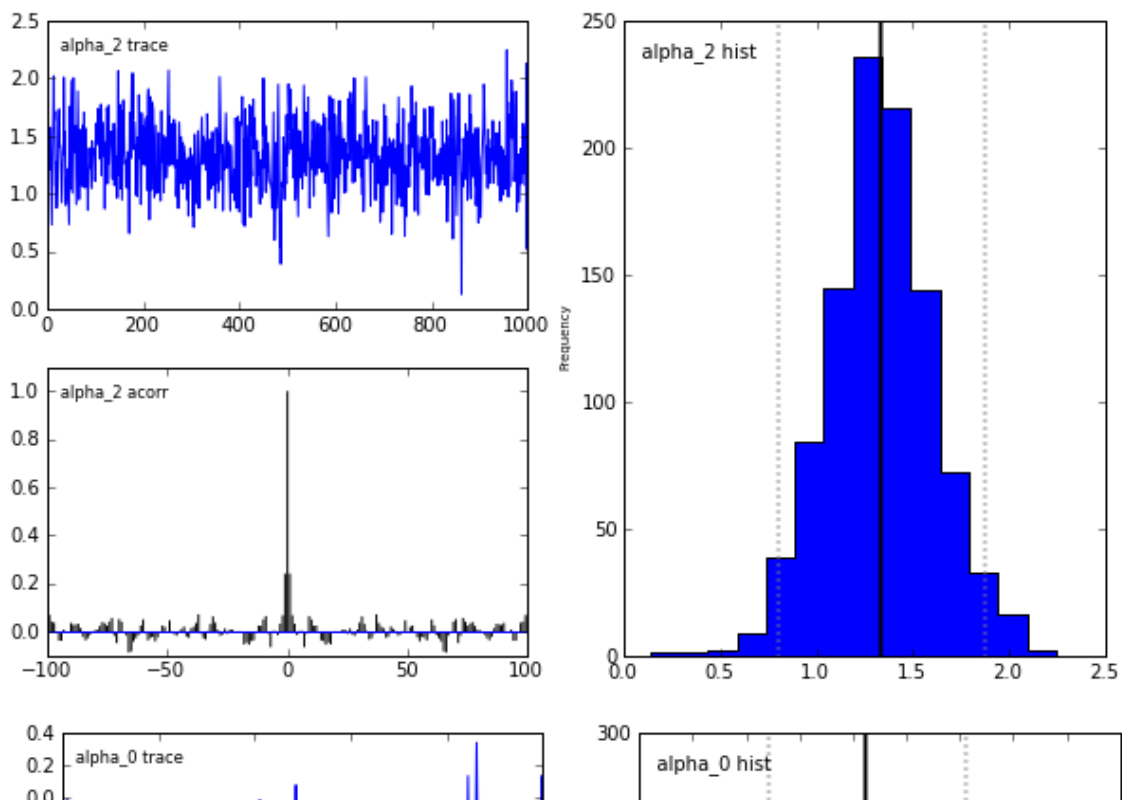
# a little longer, but not special initial values, does it converge?
%time m.sample(200000, 100000, 100, progress_bar=False)

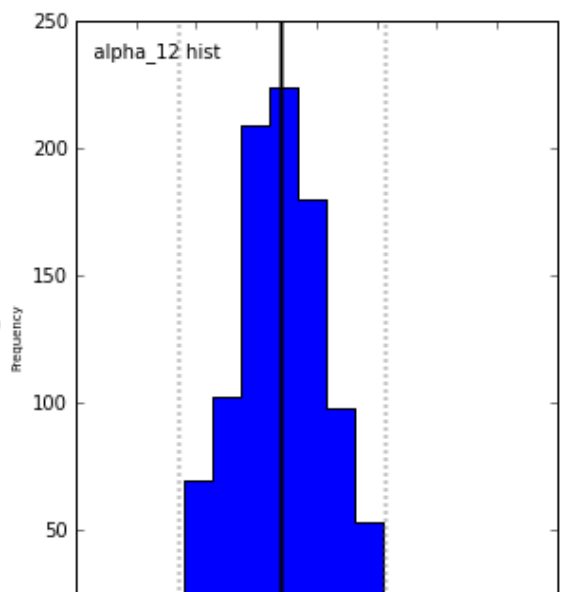
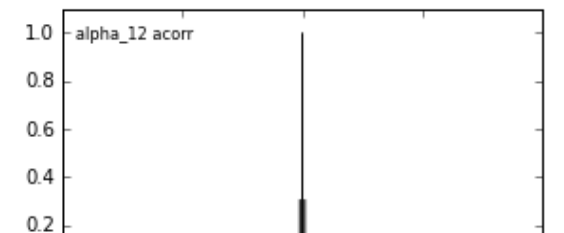
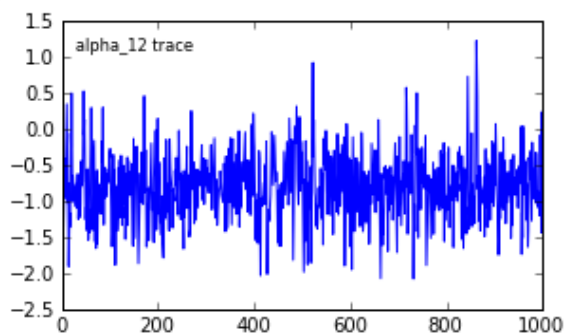
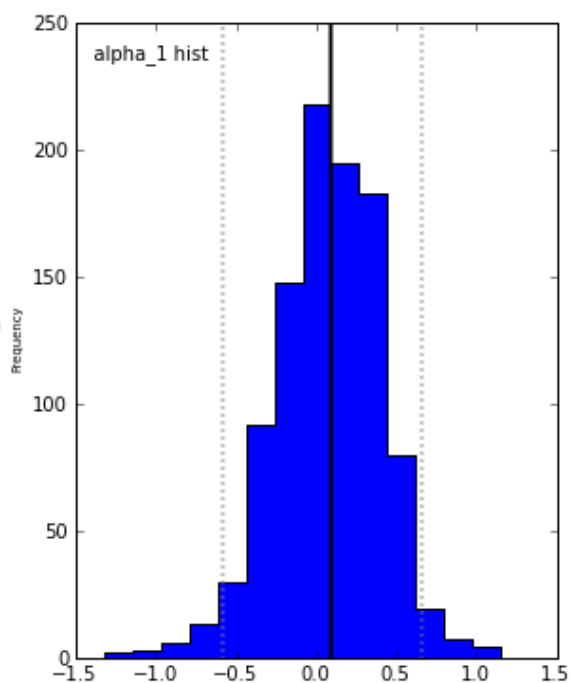
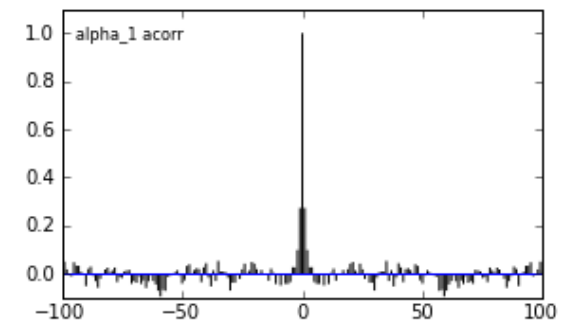
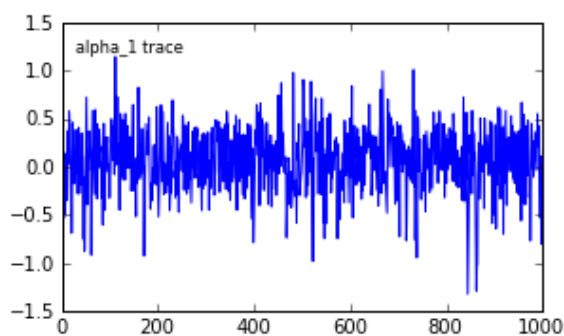
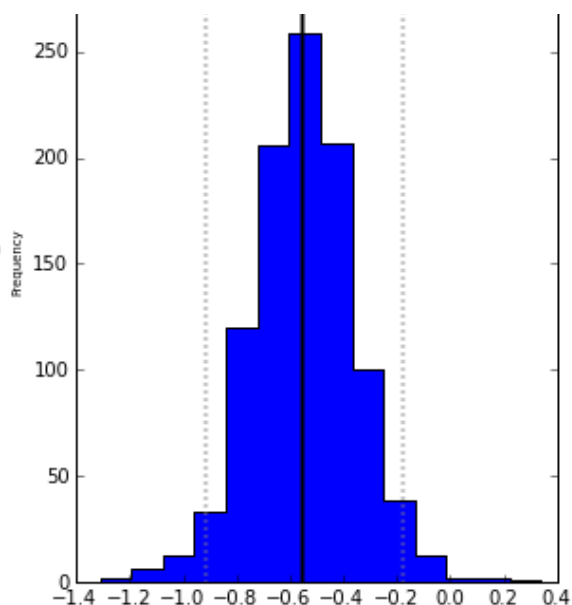
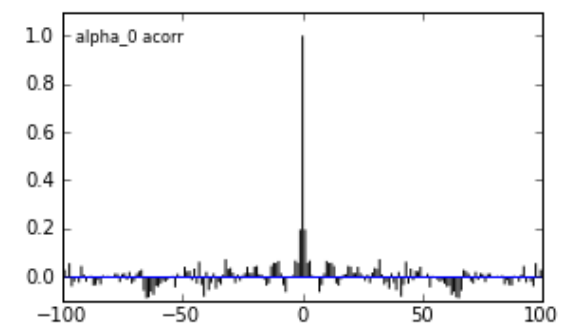
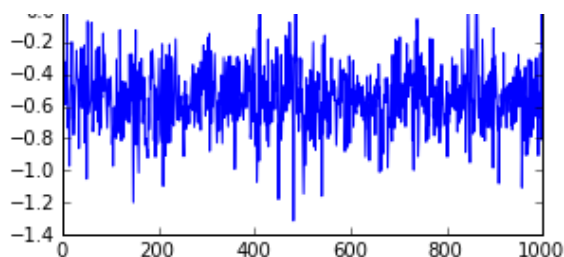
```

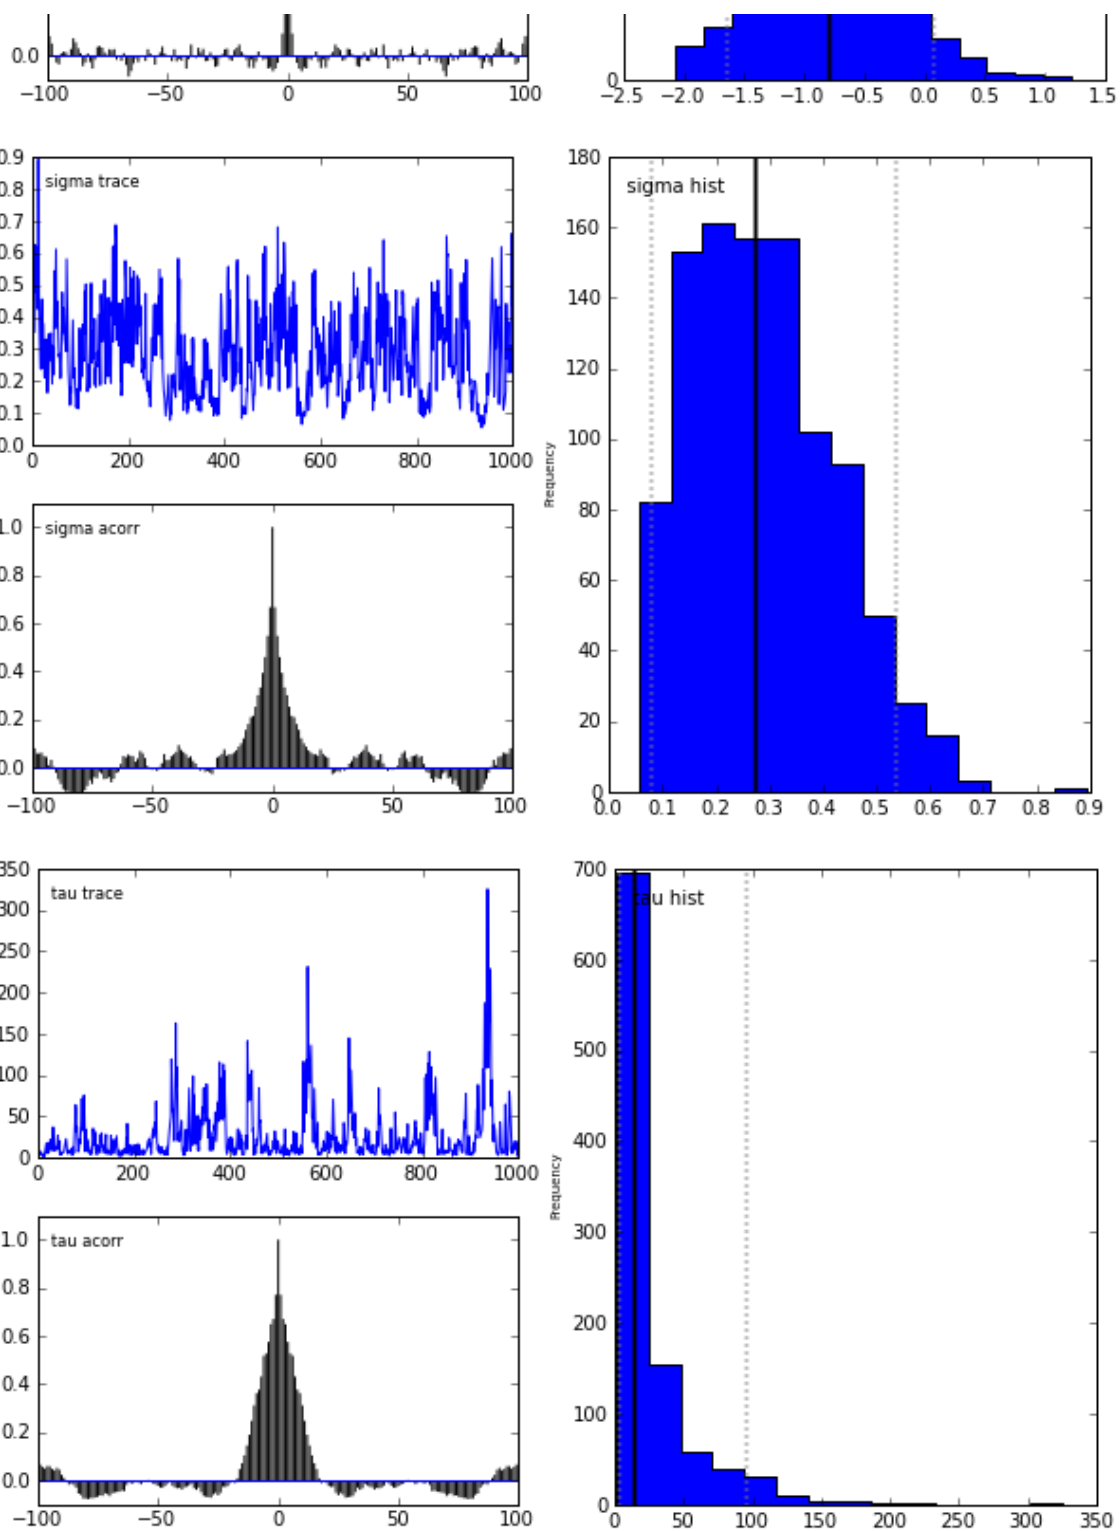
CPU times: user 123.14 s, sys: 0.64 s, total: 123.77 s
Wall time: 128.74 s

In [15]: mc.Matplot.plot(m)

Plotting alpha_2
Plotting alpha_0
Plotting alpha_1
Plotting alpha_12
Plotting sigma
Plotting tau







In []:

